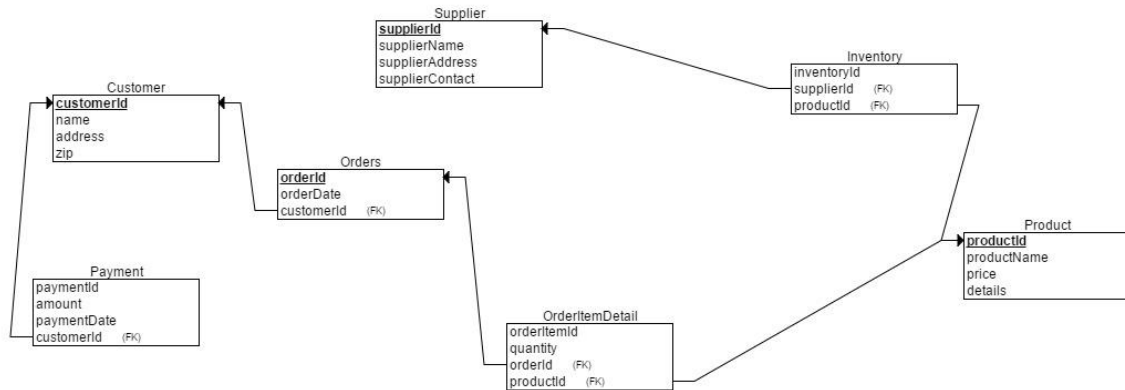


# SQLite DB ASSIGNMENT

1. Installed SQLite Add-ons successfully for Firefox
2. Designed a database for Purchase Order Management System.
3. Created a schema with necessary tables from previous step.



4. Inserted sample data successfully.
  5. Tried running different queries learnt in this chapter
- 5.1 Query to find the count of Customers

Select count(\*) as "Customer Count" from Customer ;

The screenshot shows a web-based SQLite query runner interface. At the top, there is a text input field containing the query: `Select count(*) as "Customer Count" from Customer ;`. Below the input field, there are three buttons: "Select a Query" (with a dropdown arrow), "Run SQL" (highlighted with a blue border), and "Actions" (with a dropdown arrow). To the right of these buttons, it says "Last Error: not an error". Below the buttons, the query result is displayed in a table with the title "Customer Count". The table has one row with the value "50" highlighted in green.

Customer Count
50

- 5.2 Query to list maximum product sold

Select productId , count(\*) as "Number of Orders" , sum(quantity) as "Total Quantity" from OrderItemDetails group by productId order by "Total Quantity" desc ;

Select productId , count(\*) as "Number of Orders" , sum(quantity) as "Total Quantity" from OrderItemDetails group by productId order by "Total Quantity" desc ;

Select a Query ▾	Run SQL	Actions ▾	Last Error: not an error
productId	Number of Orders	Total Quantity	
P010	3	13	
P030	3	13	
P014	3	12	
P020	3	12	
P025	2	12	
P031	2	12	
P032	3	12	
P043	4	12	
P021	2	11	
P022	3	11	
P012	3	10	
P016	3	10	
P018	3	10	
P024	3	10	
P026	3	10	
P005	2	9	

### 5.3 Query to Count of Product Supplied by each Supplier

Select supplierId ,count(\*) from Inventory group by supplierId order by count(\*) desc ;

Select a Query ▾

Run SQL

Actions ▾

Last Error: not an error

supplierId	count(*)
S08	5
S09	5
S10	5
S11	5
S12	5
S13	5
S14	5
S06	4
S07	4
S15	4
S02	3
S03	3
S04	3
S05	3

### 5.4 Query to find Suppliername ,Product details sold by supplies S08

Select S.supplierName , P.productName , P.price from Inventory I , Supplier S , Product P where I.supplierId = S.supplierId and I.productId = P.productId and I.supplierId = 'S08';

Enter SQL



Select | Data

Select S.supplierName , P.productName , P.price from Inventory I , Supplier S , Product P where  
I.supplierId = S.supplierId and I.productId = P.productId and I.supplierId = 'S08';

Select a Query ▾

Run SQL

Actions ▾

Last Error:

not an error

supplierName

productName

Specialty Biscuits, Ltd.

Konbu

Specialty Biscuits, Ltd.

Jack's New England Clam Chowder

Specialty Biscuits, Ltd.

Uncle Bob's Organic Dried Pears

Specialty Biscuits, Ltd.

Sir Rodney's Scones

Specialty Biscuits, Ltd.

Nord-Ost Matjeshering

# DB2 EXPRESS C ASSIGNMENT

1. We created database name SAMPLE with force parameter (using: db2sampl command)

Command used: write the following command in db2 command window

```
db2sampl -force
```

2. We created a SAMLE database with schema name Department which contains tables named Courses, Faculties, Students and Enrollements. Below are the structures for each table in Department schema:

db2 => Describe table enrollements

Column name	Data type schema	Data type name	Column Length	Scale	Nulls
COURSEID	SYSIBM	VARCHAR	10	0	No
STUDENTID	SYSIBM	VARCHAR	10	0	No
FACULTYID	SYSIBM	VARCHAR	10	0	No
ENROLLEMENTDATE	SYSIBM	DATE	4	0	Yes

4 record(s) selected.

db2 => Describe table courses

Column name	Data type schema	Data type name	Column Length	Scale	Nulls
COURSEID	SYSIBM	VARCHAR	10	0	No
COURSENAME	SYSIBM	VARCHAR	30	0	No
COURSEDESCRIPTION	SYSIBM	VARCHAR	100	0	Yes
CREDITS	SYSIBM	INTEGER	4	0	Yes
NUMBEROFSTUDENTS	SYSIBM	VARCHAR	5	0	Yes

5 record(s) selected.

db2 => Describe table students

Column name	Data type schema	Data type name	Column Length	Scale	Nulls
STUDENTID	SYSIBM	VARCHAR	10	0	No
STUDENTNAME	SYSIBM	VARCHAR	30	0	No
CONTACT	SYSIBM	VARCHAR	10	0	Yes

3 record(s) selected.

db2 => describe table faculties

Column name	Data type schema	Data type name	Column Length	Scale	Nulls
FACULTYID	SYSIBM	VARCHAR	10	0	No
FACULTYNAME	SYSIBM	VARCHAR	30	0	No
FACULTYEMAILID	SYSIBM	VARCHAR	20	0	Yes

3 record(s) selected.

db2 =>

3. We ran two queries (use where clause and Group by). Blow are the snapshots of sample queries we ran:

#### QUERY#1

Query to print the courseid ,coursename,facultyid,facultyname grouped by the courseid and in ascending order of there courseid :

```
select Enrollements.courseID,Courses.CourseName,Enrollements.facultyID,Faculties.facultyName,
COUNT(Courses.CourseID) "Numberof students" From enrollements JOIN courses ON
Enrollements.courseID = courses.courseID JOIN Faculties ON Faculties.FacultyID = Enrollements.FacultyID
GROUP BY Enrollements.CourseID,Courses.CourseName,Enrollements.FacultyID,Faculties.FacultyName
ORDER BY Enrollements.CourseID;
```

OUTPUT:

COURSEID	COURSENAME	FACULTYID	FACULTYNAME	Numberof students
CMPE180-38	Database	F006	Adams	2
CMPE180-92	DataStructures in C++	F007	Mayank	3
CMPE180-94	Operating System	F003	Lee	2
CMPE272	EnterpriseSoftwarePlatform	F001	George	2
CMPE272	EnterpriseSoftwarePlatform	F002	K Patel	2
CMPE273	ENTERPRISE DISTRIBUTED SYSTEM	F004	Sagar	1
CMPE273	ENTERPRISE DISTRIBUTED SYSTEM	F007	Mayank	1
CMPE281	Cloud computing	F001	George	3
CMPE283	Virtualization	F005	Ramesh	3

9 record(s) selected.

db2 =>

#### QUERY#2

Query to print coursedid,CourseName,studentID,studentName,facultyID,facultyName for all the faculty with faculty as F001 or F006 in ascending order of there CourseID.

```
SelectEnrollements.courseID,Courses.CourseName,Enrollements.studentID,Students.studentName,Enrolle
ments.facultyID,Faculties.facultyName From enrollements JOIN courses ON Enrollements.courseID =
courses.courseID JOIN students ON students.studentID = Enrollements.studentID JOIN Faculties ON
Faculties.FacultyID = Enrollements.FacultyID where Enrollements.facultyID = 'F001' OR
Enrollements.facultyID = 'F006' ORDER BY Enrollements.CourseID;
```

OUTPUT:

COURSEID	COURSENAME	STUDENTID	STUDENTNAME	FACULTYID	FACULTYNAME
CMPE180-38	Database	S001	Rahil	F006	Adams
CMPE180-38	Database	S005	A MICHAEL	F006	Adams
CMPE272	EnterpriseSoftwarePlatform	S003	SANKET	F001	George
CMPE272	EnterpriseSoftwarePlatform	S004	Parth	F001	George
CMPE281	Cloud computing	S004	Parth	F001	George
CMPE281	Cloud computing	S002	Anup	F001	George
CMPE281	Cloud computing	S005	A MICHAEL	F001	George

7 record(s) selected.

4. Generate query explain plan (use: db2exfmt tool)

Command Used:

db2 connect to SAMPLE(database name)

db2 set schema department

```
db2 set current explain mode explain
db2 -tvf Query.sql (Query file with path if not in the current directory)
db2 set current explain mode no
db2exfmt -d SAMPLE -1 -o -t(for terminal output)
db2exfmt -d sample -1 -o C:\Users\R@HIL\db2exfmt_query_explain_output.txt
```

5. Content of the Explain Plan generated in the above text file

```
DB2 Universal Database Version 11.1, 5622-044 (c) Copyright IBM
Corp. 1991, 2015
Licensed Material - Program Property of IBM
IBM DATABASE 2 Explain Table Format Tool
```

```
***** EXPLAIN INSTANCE *****
```

```
DB2_VERSION:      11.01.1
FORMATTED ON DB:  SAMPLE
SOURCE_NAME:      SQLC2026
SOURCE_SCHEMA:    NULLID
SOURCE_VERSION:
EXPLAIN_TIME:     2017-02-27-21.41.31.054000
EXPLAIN_REQUESTER: R@HIL
```

Database Context:

-----

```
Parallelism:      None
CPU Speed:        4.723442e-007
Comm Speed:       0
Buffer Pool size: 250
Sort Heap size:   256
Database Heap size: 600
Lock List size:   4096
Maximum Lock List: 22
Average Applications: 1
Locks Available:  28835
```

Package Context:

-----

```
SQL Type:         Dynamic
Optimization Level: 5
Blocking:         Block All Cursors
Isolation Level:  Cursor Stability
```

```
----- STATEMENT 1  SECTION 201 -----
```

```
QUERYNO:          2
QUERYTAG:          CLP
Statement Type:    Select
Updatable:         No
Deletable:         No
Query Degree:      1
```

Original Statement:

-----

```
select
  A.courseID,
  B.CourseName,
  A.facultyID,
  C.facultyName,
```

```

A.studentID,
D.studentName,
A.enrollementDate
From
  enrollements A JOIN courses B
ON A.courseID = B.courseID JOIN students D
ON D.studentID = A.studentID JOIN Faculties C
ON C.FacultyID = A.FacultyID
where
  (A.facultyID = 'F001' OR
   A.facultyID = 'F006' OR
   A.facultyID = 'F003') AND
  (A.studentID = 'S001' OR
   A.studentID = 'S002' OR
   A.studentID = 'S004')
ORDER BY
  A.CourseID

```

Optimized Statement:

```

-----
SELECT
  Q6.COURSEID AS "COURSEID",
  Q5.COURSENAME AS "COURSENAME",
  Q6.FACULTYID AS "FACULTYID",
  Q8.FACULTYNAME AS "FACULTYNAME",
  Q6.STUDENTID AS "STUDENTID",
  Q7.STUDENTNAME AS "STUDENTNAME",
  Q6.ENROLLEMENTDATE AS "ENROLLEMENTDATE"
FROM
  DEPARTMENT.COURSES AS Q5,
  DEPARTMENT.ENROLLEMENTS AS Q6,
  DEPARTMENT.STUDENTS AS Q7,
  DEPARTMENT.FACULTIES AS Q8
WHERE
  (Q6.COURSEID = Q5.COURSEID) AND
  (Q7.STUDENTID = Q6.STUDENTID) AND
  (Q8.FACULTYID = Q6.FACULTYID) AND
  Q6.STUDENTID IN ('S001', 'S002', 'S004') AND
  Q6.FACULTYID IN ('F001', 'F006', 'F003')
ORDER BY
  Q6.COURSEID

```

Access Plan:

```

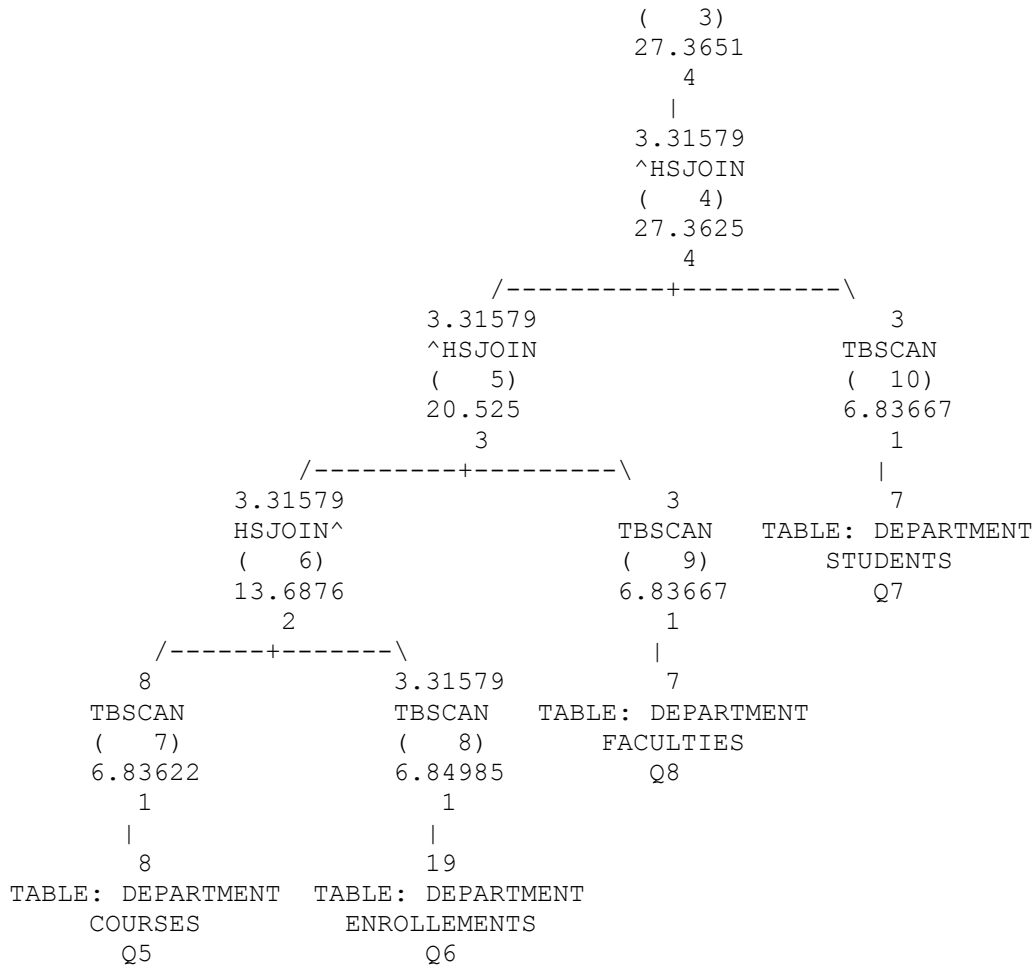
-----
          Total Cost:          27.3662
          Query Degree:         1

```

```

          Rows
RETURN
(   1)
Cost
I/O
|
3.31579
TBSCAN
(   2)
27.3662
4
|
3.31579
SORT

```



Operator Symbols :

Symbol	Description
>JOIN	: Left outer join
JOIN<	: Right outer join
>JOIN<	: Full outer join
xJOIN	: Left antijoin
JOINx	: Right antijoin
^JOIN	: Left early out
JOIN^	: Right early out

Extended Diagnostic Information:

Diagnostic Identifier: 1

Diagnostic Details: EXP0256I Analysis of the query shows that the query might execute faster if an additional index was created to enable zigzag join. Schema name: "DEPARTMENT". Table name: "ENROLLEMENTS". Column list: "(STUDENTID, FACULTYID, COURSEID)".

Plan Details:

1) RETURN: (Return Result)



Cumulative Total Cost: 27.3662  
 Cumulative CPU Cost: 309521  
 Cumulative I/O Cost: 4  
 Cumulative Re-Total Cost: 27.3636  
 Cumulative Re-CPU Cost: 303956  
 Cumulative Re-I/O Cost: 0  
 Cumulative First Row Cost: 27.3658  
 Estimated Bufferpool Buffers: 0

Arguments:

-----  
 BLDLEVEL: (Build level)  
           DB2 v11.1.1010.160 : s1612051900  
 HEAPUSE : (Maximum Statement Heap Usage)  
           208 Pages  
 PLANID : (Access plan identifier)  
           44939729a14955e7  
 PREPTIME: (Statement prepare time)  
           218 milliseconds  
 SEMEVID : (Semantic environment identifier)  
           4465607c101ab241  
 STMTHEAP: (Statement heap size)  
           8192  
 STMTID : (Normalized statement identifier)  
           8105e04819c65e82

Input Streams:

-----  
           13) From Operator #2  
  
                     Estimated number of rows: 3.31579  
                     Number of columns: 7  
                     Subquery predicate ID: Not

Applicable

Column Names:

-----  
 +Q9.COURSEID(A)+Q9.ENROLLEMENTDATE  
 +Q9.STUDENTNAME+Q9.STUDENTID+Q9.FACULTYNAME  
 +Q9.FACULTYID+Q9.COURSENAME

2) TBSCAN: (Table Scan)

Cumulative Total Cost: 27.3662  
 Cumulative CPU Cost: 309521  
 Cumulative I/O Cost: 4  
 Cumulative Re-Total Cost: 27.3636  
 Cumulative Re-CPU Cost: 303956  
 Cumulative Re-I/O Cost: 0  
 Cumulative First Row Cost: 27.3658  
 Estimated Bufferpool Buffers: 0

Arguments:

-----  
 MAXPAGES: (Maximum pages for prefetch)  
           ALL  
 PREFETCH: (Type of Prefetch)  
           NONE  
 SCANDIR : (Scan Direction)  
           FORWARD  
 SPEED : (Assumed speed of scan, in sharing structures)  
           SLOW  
 THROTTLE: (Scan may be throttled, for scan sharing)

FALSE  
VISIBLE : (May be included in scan sharing structures)  
FALSE  
WRAPPING: (Scan may start anywhere and wrap)  
FALSE

Input Streams:

-----

12) From Operator #3

Estimated number of rows: 3.31579  
Number of columns: 7  
Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q6.COURSEID(A)+Q5.COURSENAME  
+Q6.ENROLLEMENTDATE+Q6.FACULTYID+Q6.STUDENTID  
+Q7.STUDENTNAME+Q8.FACULTYNAME

Output Streams:

-----

13) To Operator #1

Estimated number of rows: 3.31579  
Number of columns: 7  
Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q9.COURSEID(A)+Q9.ENROLLEMENTDATE  
+Q9.STUDENTNAME+Q9.STUDENTID+Q9.FACULTYNAME  
+Q9.FACULTYID+Q9.COURSENAME

3) SORT : (Sort)

Cumulative Total Cost: 27.3651  
Cumulative CPU Cost: 307155  
Cumulative I/O Cost: 4  
Cumulative Re-Total Cost: 27.3625  
Cumulative Re-CPU Cost: 301590  
Cumulative Re-I/O Cost: 0  
Cumulative First Row Cost: 27.3651  
Estimated Bufferpool Buffers: 1

Arguments:

-----

DUPLWARN: (Duplicates Warning flag)  
FALSE  
KEYS : (Key cardinality)  
4  
NUMROWS : (Estimated number of rows)  
4  
ROWWIDTH: (Estimated width of rows)  
80.000000  
SORTKEY : (Sort Key column)  
1: Q6.COURSEID(A)  
TEMPSIZE: (Temporary Table Page Size)  
8192  
UNIQUE : (Uniqueness required flag)

FALSE

Input Streams:

-----

11) From Operator #4

Estimated number of rows:	3.31579	
Number of columns:	8	
Subquery predicate ID:		Not

Applicable

Column Names:

-----

+Q5.COURSENAME+Q6.ENROLLEMENTDATE+Q6.FACULTYID  
+Q6.STUDENTID+Q6.COURSEID+Q7.STUDENTNAME  
+Q7.STUDENTID+Q8.FACULTYNAME

Output Streams:

-----

12) To Operator #2

Estimated number of rows:	3.31579	
Number of columns:	7	
Subquery predicate ID:		Not

Applicable

Column Names:

-----

+Q6.COURSEID (A) +Q5.COURSENAME  
+Q6.ENROLLEMENTDATE+Q6.FACULTYID+Q6.STUDENTID  
+Q7.STUDENTNAME+Q8.FACULTYNAME

4) HSJOIN: (Hash Join)

Cumulative Total Cost:	27.3625
Cumulative CPU Cost:	301590
Cumulative I/O Cost:	4
Cumulative Re-Total Cost:	27.3625
Cumulative Re-CPU Cost:	301590
Cumulative Re-I/O Cost:	4
Cumulative First Row Cost:	27.3625
Estimated Bufferpool Buffers:	1

Arguments:

-----

BITFLTR : (Hash Join Bit Filter used)  
FALSE  
EARLYOUT: (Early Out flag)  
LEFT  
HASHCODE: (Hash Code Size)  
24 BIT  
HASHTBSZ: (Number of hash table entries)  
3  
TEMPSIZE: (Temporary Table Page Size)  
8192  
TUPBLKSZ: (Tuple Block Size (bytes))  
4000

Predicates:

-----

7) Predicate used in Join,  
Comparison Operator: Equal (=)

Subquery Input Required: No  
Filter Factor: 0.333333

Predicate Text:  
-----  
(Q7.STUDENTID = Q6.STUDENTID)

Input Streams:  
-----

8) From Operator #5

Estimated number of rows: 3.31579  
Number of columns: 7  
Subquery predicate ID: Not

Applicable

Column Names:  
-----  
+Q5.COURSENAME+Q6.ENROLLEMENTDATE+Q6.FACULTYID  
+Q6.STUDENTID+Q6.COURSEID+Q8.FACULTYNAME  
+Q8.FACULTYID

10) From Operator #10

Estimated number of rows: 3  
Number of columns: 2  
Subquery predicate ID: Not

Applicable

Column Names:  
-----  
+Q7.STUDENTNAME+Q7.STUDENTID

Output Streams:  
-----

11) To Operator #3

Estimated number of rows: 3.31579  
Number of columns: 8  
Subquery predicate ID: Not

Applicable

Column Names:  
-----  
+Q5.COURSENAME+Q6.ENROLLEMENTDATE+Q6.FACULTYID  
+Q6.STUDENTID+Q6.COURSEID+Q7.STUDENTNAME  
+Q7.STUDENTID+Q8.FACULTYNAME

5) HSJOIN: (Hash Join)

Cumulative Total Cost: 20.525  
Cumulative CPU Cost: 232945  
Cumulative I/O Cost: 3  
Cumulative Re-Total Cost: 20.525  
Cumulative Re-CPU Cost: 232945  
Cumulative Re-I/O Cost: 3  
Cumulative First Row Cost: 20.525  
Estimated Bufferpool Buffers: 1

Arguments:

-----  
BITFLTR : (Hash Join Bit Filter used)  
FALSE  
EARLYOUT: (Early Out flag)  
LEFT  
HASHCODE: (Hash Code Size)  
24 BIT  
HASHTBSZ: (Number of hash table entries)  
3  
JN INPUT: (Join input leg)  
OUTER  
TEMPSIZE: (Temporary Table Page Size)  
8192  
TUPBLKSZ: (Tuple Block Size (bytes))  
4000

Predicates:

-----  
8) Predicate used in Join,  
Comparison Operator: Equal (=)  
Subquery Input Required: No  
Filter Factor: 0.333333

Predicate Text:  
-----  
(Q8.FACULTYID = Q6.FACULTYID)

Input Streams:

-----  
5) From Operator #6  
  
Estimated number of rows: 3.31579  
Number of columns: 6  
Subquery predicate ID: Not  
  
Column Names:  
-----  
+Q5.COURSENAME+Q5.COURSEID+Q6.ENROLLEMENTDATE  
+Q6.FACULTYID+Q6.STUDENTID+Q6.COURSEID

Applicable

7) From Operator #9  
  
Estimated number of rows: 3  
Number of columns: 2  
Subquery predicate ID: Not  
  
Column Names:  
-----  
+Q8.FACULTYNAME+Q8.FACULTYID

Applicable

Output Streams:

-----  
8) To Operator #4  
  
Estimated number of rows: 3.31579  
Number of columns: 7  
Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q5.COURSENAME+Q6.ENROLLEMENTDATE+Q6.FACULTYID  
+Q6.STUDENTID+Q6.COURSEID+Q8.FACULTYNAME  
+Q8.FACULTYID

6) HSJOIN: (Hash Join)

Cumulative Total Cost: 13.6876  
Cumulative CPU Cost: 164303  
Cumulative I/O Cost: 2  
Cumulative Re-Total Cost: 13.6876  
Cumulative Re-CPU Cost: 164303  
Cumulative Re-I/O Cost: 2  
Cumulative First Row Cost: 13.6876  
Estimated Bufferpool Buffers: 1

Arguments:

-----

BITFLTR : (Hash Join Bit Filter used)  
FALSE  
EARLYOUT: (Early Out flag)  
RIGHT  
HASHCODE: (Hash Code Size)  
24 BIT  
HASHTBSZ: (Number of hash table entries)  
3  
JN INPUT: (Join input leg)  
OUTER  
TEMPSIZE: (Temporary Table Page Size)  
8192  
TUPBLKSZ: (Tuple Block Size (bytes))  
4000

Predicates:

-----

6) Predicate used in Join,  
Comparison Operator: Equal (=)  
Subquery Input Required: No  
Filter Factor: 0.125

Predicate Text:

-----

(Q6.COURSEID = Q5.COURSEID)

Input Streams:

-----

2) From Operator #7

Estimated number of rows: 8  
Number of columns: 2  
Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q5.COURSENAME+Q5.COURSEID

4) From Operator #8

	Estimated number of rows:	3.31579	
	Number of columns:	4	
Applicable	Subquery predicate ID:		Not

Column Names:

-----

+Q6.ENROLLEMENTDATE+Q6.FACULTYID+Q6.STUDENTID  
+Q6.COURSEID

Output Streams:

-----

5) To Operator #5

	Estimated number of rows:	3.31579	
	Number of columns:	6	
Applicable	Subquery predicate ID:		Not

Column Names:

-----

+Q5.COURSENAME+Q5.COURSEID+Q6.ENROLLEMENTDATE  
+Q6.FACULTYID+Q6.STUDENTID+Q6.COURSEID

7) TBSCAN: (Table Scan)

Cumulative Total Cost:	6.83622
Cumulative CPU Cost:	66091
Cumulative I/O Cost:	1
Cumulative Re-Total Cost:	0.00732039
Cumulative Re-CPU Cost:	15498
Cumulative Re-I/O Cost:	0
Cumulative First Row Cost:	6.82971
Estimated Bufferpool Buffers:	1

Arguments:

-----

CUR\_COMM: (Currently Committed)  
TRUE

JN INPUT: (Join input leg)  
OUTER

LCKAVOID: (Lock Avoidance)  
TRUE

MAXPAGES: (Maximum pages for prefetch)  
ALL

PREFETCH: (Type of Prefetch)  
NONE

ROWLOCK : (Row Lock intent)  
SHARE (CS/RS)

SCANDIR : (Scan Direction)  
FORWARD

SKIP\_INS: (Skip Inserted Rows)  
TRUE

SPEED : (Assumed speed of scan, in sharing structures)  
FAST

TABLOCK : (Table Lock intent)  
INTENT SHARE

TBISOLVL: (Table access Isolation Level)  
CURSOR STABILITY

THROTTLE: (Scan may be throttled, for scan sharing)  
TRUE

VISIBLE : (May be included in scan sharing structures)

TRUE  
WRAPPING: (Scan may start anywhere and wrap)  
TRUE

Input Streams:

-----

1) From Object DEPARTMENT.COURSES

Estimated number of rows:	8	
Number of columns:	3	
Subquery predicate ID:		Not

Applicable

Column Names:

-----

+Q5.\$RID\$+Q5.COURSENAME+Q5.COURSEID

Output Streams:

-----

2) To Operator #6

Estimated number of rows:	8	
Number of columns:	2	
Subquery predicate ID:		Not

Applicable

Column Names:

-----

+Q5.COURSENAME+Q5.COURSEID

8) TBSCAN: (Table Scan)

Cumulative Total Cost:	6.84985
Cumulative CPU Cost:	94952
Cumulative I/O Cost:	1
Cumulative Re-Total Cost:	0.0209527
Cumulative Re-CPU Cost:	44359
Cumulative Re-I/O Cost:	0
Cumulative First Row Cost:	6.83491
Estimated Bufferpool Buffers:	1

Arguments:

-----

CUR\_COMM: (Currently Committed)

TRUE

JN INPUT: (Join input leg)

INNER

LCKAVOID: (Lock Avoidance)

TRUE

MAXPAGES: (Maximum pages for prefetch)

ALL

PREFETCH: (Type of Prefetch)

NONE

ROWLOCK : (Row Lock intent)

SHARE (CS/RS)

SCANDIR : (Scan Direction)

FORWARD

SKIP\_INS: (Skip Inserted Rows)

TRUE

SPEED : (Assumed speed of scan, in sharing structures)

FAST

TABLOCK : (Table Lock intent)



INTENT SHARE  
TBISOLVL: (Table access Isolation Level)  
CURSOR STABILITY  
THROTTLE: (Scan may be throttled, for scan sharing)  
TRUE  
VISIBLE : (May be included in scan sharing structures)  
TRUE  
WRAPPING: (Scan may start anywhere and wrap)  
TRUE

Predicates:

-----

11) Sargable Predicate,  
Comparison Operator: In List (IN), evaluated by  
binary search (list sorted at compile-time)  
Subquery Input Required: No  
Filter Factor: 0.368421

Predicate Text:

-----

Q6.STUDENTID IN ('S001', 'S002', 'S004')

12) Sargable Predicate,  
Comparison Operator: In List (IN), evaluated by  
binary search (list sorted at compile-time)  
Subquery Input Required: No  
Filter Factor: 0.473684

Predicate Text:

-----

Q6.FACULTYID IN ('F001', 'F006', 'F003')

Input Streams:

-----

3) From Object DEPARTMENT.ENROLLEMENTS

Estimated number of rows: 19  
Number of columns: 5  
Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q6.\$RID\$+Q6.ENROLLEMENTDATE+Q6.FACULTYID  
+Q6.STUDENTID+Q6.COURSEID

Output Streams:

-----

4) To Operator #6

Estimated number of rows: 3.31579  
Number of columns: 4  
Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q6.ENROLLEMENTDATE+Q6.FACULTYID+Q6.STUDENTID  
+Q6.COURSEID

9) TBSCAN: (Table Scan)

Cumulative Total Cost:	6.83667
Cumulative CPU Cost:	67039.5
Cumulative I/O Cost:	1
Cumulative Re-Total Cost:	0.00776841
Cumulative Re-CPU Cost:	16446.5
Cumulative Re-I/O Cost:	0
Cumulative First Row Cost:	6.83118
Estimated Bufferpool Buffers:	1

Arguments:

-----

CUR\_COMM: (Currently Committed)  
TRUE  
JN INPUT: (Join input leg)  
INNER  
LCKAVOID: (Lock Avoidance)  
TRUE  
MAXPAGES: (Maximum pages for prefetch)  
ALL  
PREFETCH: (Type of Prefetch)  
NONE  
ROWLOCK : (Row Lock intent)  
SHARE (CS/RS)  
SCANDIR : (Scan Direction)  
FORWARD  
SKIP\_INS: (Skip Inserted Rows)  
TRUE  
SPEED : (Assumed speed of scan, in sharing structures)  
FAST  
TABLOCK : (Table Lock intent)  
INTENT SHARE  
TBISOLVL: (Table access Isolation Level)  
CURSOR STABILITY  
THROTTLE: (Scan may be throttled, for scan sharing)  
TRUE  
VISIBLE : (May be included in scan sharing structures)  
TRUE  
WRAPPING: (Scan may start anywhere and wrap)  
TRUE

Predicates:

-----

9) Sargable Predicate,  
Comparison Operator: In List (IN), evaluated by  
binary search (list sorted at compile-time)  
Subquery Input Required: No  
Filter Factor: 0.428571

Predicate Text:

-----

Q8.FACULTYID IN ('F001', 'F006', 'F003')

Input Streams:

-----

6) From Object DEPARTMENT.FACULTIES

Estimated number of rows:	7
Number of columns:	3

Applicable Subquery predicate ID: Not

Column Names:

-----

+Q8.\$RID\$+Q8.FACULTYNAME+Q8.FACULTYID

Output Streams:

-----

7) To Operator #5

Estimated number of rows: 3

Number of columns: 2

Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q8.FACULTYNAME+Q8.FACULTYID

10) TBSCAN: (Table Scan)

Cumulative Total Cost: 6.83667

Cumulative CPU Cost: 67039.5

Cumulative I/O Cost: 1

Cumulative Re-Total Cost: 0.00776841

Cumulative Re-CPU Cost: 16446.5

Cumulative Re-I/O Cost: 0

Cumulative First Row Cost: 6.83118

Estimated Bufferpool Buffers: 1

Arguments:

-----

CUR\_COMM: (Currently Committed)

TRUE

JN INPUT: (Join input leg)

INNER

LCKAVOID: (Lock Avoidance)

TRUE

MAXPAGES: (Maximum pages for prefetch)

ALL

PREFETCH: (Type of Prefetch)

NONE

ROWLOCK : (Row Lock intent)

SHARE (CS/RS)

SCANDIR : (Scan Direction)

FORWARD

SKIP\_INS: (Skip Inserted Rows)

TRUE

SPEED : (Assumed speed of scan, in sharing structures)

FAST

TABLOCK : (Table Lock intent)

INTENT SHARE

TBISOLVL: (Table access Isolation Level)

CURSOR STABILITY

THROTTLE: (Scan may be throttled, for scan sharing)

TRUE

VISIBLE : (May be included in scan sharing structures)

TRUE

WRAPPING: (Scan may start anywhere and wrap)

TRUE

Predicates:

-----

10) Sargable Predicate,  
Comparison Operator: In List (IN), evaluated by  
binary search (list sorted at compile-time)  
Subquery Input Required: No  
Filter Factor: 0.428571

Predicate Text:

-----

Q7.STUDENTID IN ('S001', 'S002', 'S004')

Input Streams:

-----

9) From Object DEPARTMENT.STUDENTS

Estimated number of rows: 7  
Number of columns: 3  
Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q7.\$RID\$+Q7.STUDENTNAME+Q7.STUDENTID

Output Streams:

-----

10) To Operator #4

Estimated number of rows: 3  
Number of columns: 2  
Subquery predicate ID: Not

Applicable

Column Names:

-----

+Q7.STUDENTNAME+Q7.STUDENTID

Objects Used in Access Plan:

-----

Schema: DEPARTMENT

Name: COURSES

Type: Table

Time of creation: 2017-02-20-00.07.05.895001  
Last statistics update: 2017-02-20-11.28.10.379000  
Number of columns: 5  
Number of rows: 8  
Width of rows: 66  
Number of buffer pool pages: 1  
Number of data partitions: 1  
Distinct row values: No  
Tablespace name: IBMDB2SAMPLEREL  
Tablespace overhead: 6.725000  
Tablespace transfer rate: 0.080000  
Source for statistics: Single Node  
Prefetch page count: 32  
Container extent page count: 32  
Table overflow record count: 0

Table Active Blocks:	-1	
Average Row Compression Ratio:		0
Percentage Rows Compressed:	0	
Average Compressed Row Size:	0	

Schema: DEPARTMENT  
 Name: ENROLLEMENTS  
 Type: Table

Time of creation:	2017-02-20-12.21.04.415001
Last statistics update:	2017-02-20-21.53.12.213000
Number of columns:	4
Number of rows:	19
Width of rows:	39
Number of buffer pool pages:	1
Number of data partitions:	1
Distinct row values:	No
Tablespace name:	IBMDB2SAMPLEREL
Tablespace overhead:	6.725000
Tablespace transfer rate:	0.080000
Source for statistics:	Single Node
Prefetch page count:	32
Container extent page count:	32
Table overflow record count:	0
Table Active Blocks:	-1
Average Row Compression Ratio:	0
Percentage Rows Compressed:	0
Average Compressed Row Size:	0

Schema: DEPARTMENT  
 Name: FACULTIES  
 Type: Table

Time of creation:	2017-02-20-11.27.45.674000
Last statistics update:	2017-02-20-21.22.47.596000
Number of columns:	3
Number of rows:	7
Width of rows:	31
Number of buffer pool pages:	1
Number of data partitions:	1
Distinct row values:	No
Tablespace name:	IBMDB2SAMPLEREL
Tablespace overhead:	6.725000
Tablespace transfer rate:	0.080000
Source for statistics:	Single Node
Prefetch page count:	32
Container extent page count:	32
Table overflow record count:	0
Table Active Blocks:	-1
Average Row Compression Ratio:	0
Percentage Rows Compressed:	0
Average Compressed Row Size:	0

Schema: DEPARTMENT  
 Name: STUDENTS  
 Type: Table

Time of creation:	2017-02-20-11.44.55.568000
Last statistics update:	2017-02-20-12.09.49.582000
Number of columns:	3
Number of rows:	7
Width of rows:	32
Number of buffer pool pages:	1
Number of data partitions:	1
Distinct row values:	No
Tablespace name:	IBMDB2SAMPLEREL

Tablespace overhead:	6.725000	
Tablespace transfer rate:	0.080000	
Source for statistics:		Single Node
Prefetch page count:	32	
Container extent page count:	32	
Table overflow record count:	0	
Table Active Blocks:	-1	
Average Row Compression Ratio:		0
Percentage Rows Compressed:	0	
Average Compressed Row Size:	0	

# IBM GRAPH DB ASSIGNMENT

- 1.) Create an account on IBM Bluemix, navigate through "Data and analytics" section and go on Graph as a service. Through the Graph Dashboard, capture the Service credentials from the left navigation Bar.
- 2.) We are now going to authenticate each of our requests made to our graph database by storing our credentials in temporary token and using it for authentication every time a request is made.

Commands used for Authentication:

```
CREDS='{
  "credentials": {
    "apiURL": "https://ibmgraph-alpha.ng.bluemix.net/a261eac3-6956-4185-a789-8e24c475e89b/g",
    "username": "575fd2b0-c011-4b94-82cf-43244b44f3a8",
    "password": "aca0d011-26f5-485d-a446-182750278a27"
  }
}'

USER=$(echo $CREDS | jq -r '.credentials.username')
PASS=$(echo $CREDS | jq -r '.credentials.password')
URL=$(echo $CREDS | jq -r '.credentials.apiURL' | sed -E 's/(.*)\./.*\/1/' ) # remove the graph name from the apiURL
alias curl='curl --max-time 60 --connect-timeout 5 --silent --show-error' # set some defaults for curl
TOKEN=$(curl "${URL}/_session" -u "$USER:$PASS" | jq -r '["gds-token"]')
echo "Your session token is $TOKEN"
```

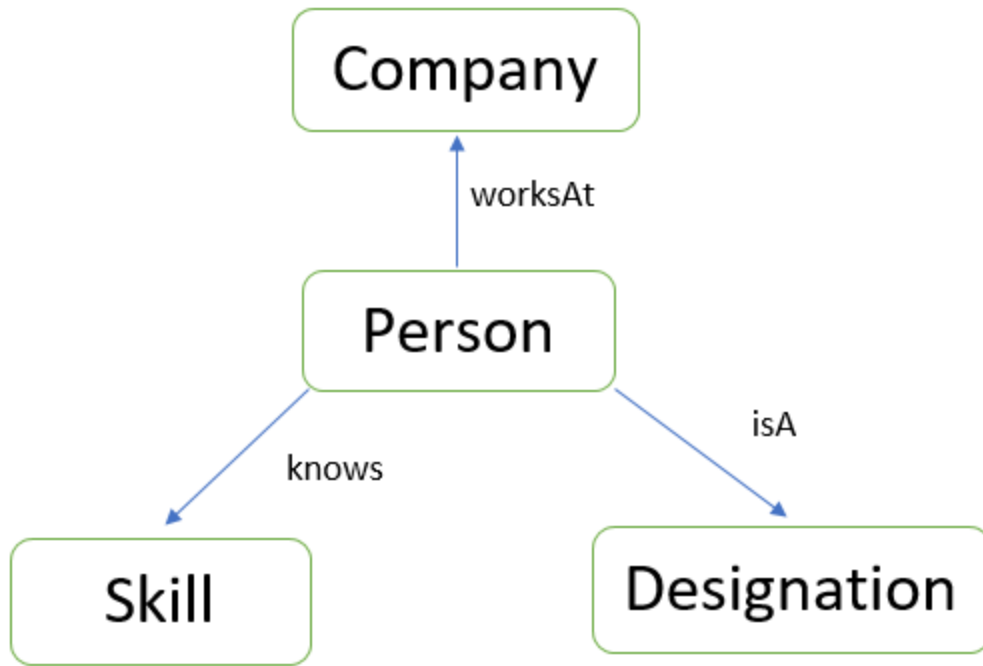
- 3.) We created a graph with the name cmpe272gp23

Commands used

```
GRAPH="cmpe272gp23"
curl "$URL/_graphs/$GRAPH" \
  -X POST \
  -H "Authorization: gds-token $TOKEN" \
  -d " | jq '.'
```

Now we are defining the schema for the graphdb cmpe272gp23

Below is the graph representation of cmpe272gp23



SCHEMA='

```

{
  "vertexLabels" : [
    {"name" : "Company"},
    {"name" : "Person"},
    {"name" : "Skill"},
    {"name" : "Position"}
  ],

  "edgeLabels": [
    {"name" : "worksAt", "multiplicity": "MANY2ONE"},
    {"name": "knows", "multiplicity" : "MULTI"},
    {"name" : "isA", "multiplicity" : "SIMPLE"}
  ],

  "propertyKeys" : [
    {"name" : "companyName", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "established", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "firstName", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "lastName", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "age", "dataType": "Integer", "cardinality": "SINGLE"},
    {"name" : "salary", "dataType": "Float", "cardinality": "SINGLE"},
    {"name" : "address", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "gender", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "email", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "skillName", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "designation", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "joiningDate", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "employeeId", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "competencyLevel", "dataType": "String", "cardinality": "SINGLE"},
    {"name" : "since", "dataType": "String", "cardinality": "SINGLE"}
  ],

  "vertexIndexes" : [
    {"name": "userbyname", "propertyKeys" : ["firstName", "lastName"], "composite": false, "unique": false },
  ]
}

```



```

        {"name":"userbyage" ,"propertyKeys" :["age"],"composite":false , "unique":false },
        {"name":"userbygender" ,"propertyKeys" :["gender"],"composite":false , "unique":false },
        {"name":"userbyskill" ,"propertyKeys" :["skillName"],"composite":false , "unique":false },
        {"name":"userbyemployeeid" ,"propertyKeys" :["employeeid"],"composite":true , "unique": true },
        {"name":"userbycompanyName" ,"propertyKeys" :["companyName"],"composite":false , "unique": false},
        {"name":"userbydesignation" ,"propertyKeys" :["designation"],"composite":false , "unique":false },
        {"name":"userbysalary" ,"propertyKeys" :["salary"],"composite":false , "unique":false }

    ],

    "edgeIndexes" : [
        {"name":"userbyjoiningDate" ,"propertyKeys" :["joiningDate"],"composite":false , "unique":false },
        {"name":"userbysince" ,"propertyKeys" :["since"],"composite":false , "unique":false },
        {"name":"userbycompetency" ,"propertyKeys" :["competencyLevel"],"composite":false , "unique":false }
    ]
}'

```

```

curl "$URL/$GRAPH/schema" \
-X POST \
-H "Authorization: gds-token $TOKEN" \
-H 'Content-Type: application/json' \
-d "$SCHEMA" | jq '.'

```

#### 4.) We are now going to insert all the data into our schema.

Commands used:

```

cat << ENDGREMLIN >gremlin.json # write everything until ENDGREMLIN into gremlin.json
{ "gremlin": "

```

```

def David = graph.addVertex(T.label, 'Person', 'firstName', 'David',
'lastName','Morgan','age',35,'salary',70000.00,'address','San
Jose','gender','Male','email','xx@gmail.com','employeeid','G001');
def Vinayak = graph.addVertex(T.label, 'Person', 'firstName', 'Vinayak',
'lastName','Patel','age',26,'salary',80000.00,'address','San
Francisco','gender','Male','email','Vinayak@gmail.com','employeeid','G002');
def Akhilesh = graph.addVertex(T.label, 'Person', 'firstName', 'Akhilesh',
'lastName','Doe','age',24,'salary',50000.00,'address','San
fernando','gender','Male','email','akhilesh@gmail.com','employeeid','G003');
def Rahil = graph.addVertex(T.label, 'Person', 'firstName', 'Rahil',
'lastName','Modi','age',30,'salary',30000.00,'address','Santa
Cruiz','gender','Male','email','Rahil@gmail.com','employeeid','G005');
def Siddharth = graph.addVertex(T.label, 'Person', 'firstName', 'Siddharth',
'lastName','Mewada','age',24,'salary',20000.00,'address','Santa
Cruiz','gender','Male','email','sid@gmail.com','employeeid','G006');

def Google = graph.addVertex(T.label, 'Company', 'companyName', 'Google', 'established', '1995');
def Apple = graph.addVertex(T.label, 'Company', 'companyName', 'Apple', 'established', '1997');

def Manager = graph.addVertex(T.label, 'Position', 'positionName','Project Manager');
def Director = graph.addVertex(T.label, 'Position', 'positionName','Director');
def Lead = graph.addVertex(T.label, 'Position', 'positionName','Technical lead');
def SystemsEngineer = graph.addVertex(T.label, 'Position', 'positionName','SystemsEngineer');

def Python = graph.addVertex(T.label, 'Skill', 'skillName', 'Python');
def Java = graph.addVertex(T.label, 'Skill', 'skillName', 'Java');
def Mainframe = graph.addVertex(T.label, 'Skill', 'skillName', 'Mainframe');

```

```
def Javascript = graph.addVertex(T.label, 'Skill', 'skillName', 'Javascript');
def Go = graph.addVertex(T.label, 'Skill', 'skillName', 'Go');
```

```
David.addEdge('knows',Python,'competencyLevel','good');
Vinayak.addEdge('knows',Java,'competencyLevel','poor');
Akhilesh.addEdge('knows',Mainframe,'competencyLevel','expert');
Rahil.addEdge('knows',Javascript,'competencyLevel','good');
Siddharth.addEdge('knows',Go,'competencyLevel','poor');
Siddharth.addEdge('knows',Python,'competencyLevel','expert');
Vinayak.addEdge('knows',Javascript,'competencyLevel','expert');
```

```
David.addEdge('worksAt',Google,'joiningDate','12 Mar 2015');
Vinayak.addEdge('worksAt',Google,'joiningDate','29 Feb 2012');
Siddharth.addEdge('worksAt',Google,'joiningDate','31 Oct 2015');
Rahil.addEdge('worksAt',Apple,'joiningDate','20 June 2013');
Akhilesh.addEdge('worksAt',Apple,'joiningDate','01 Sep 2014');
```

```
David.addEdge('isA',Manager,'since','Mar 2015');
Vinayak.addEdge('isA',SystemsEngineer,'since','Jun 2014');
Akhilesh.addEdge('isA',Director,'since','Dec 2014');
Rahil.addEdge('isA',SystemsEngineer,'since','Aug 2013');
Siddharth.addEdge('isA',Lead,'since','Nov 2015');
```

```
"
```

```
}
```

```
ENDGREMLIN
```

```
curl "$URL/$GRAPH/gremlin" \
-X POST \
-H "Authorization: gds-token $TOKEN" \
-H 'Content-Type: application/json' \
-d @gremlin.json | jq '.'
```

5.) Now we are going to use Gremlin queries to traverse through our graph database cmpe272gp23

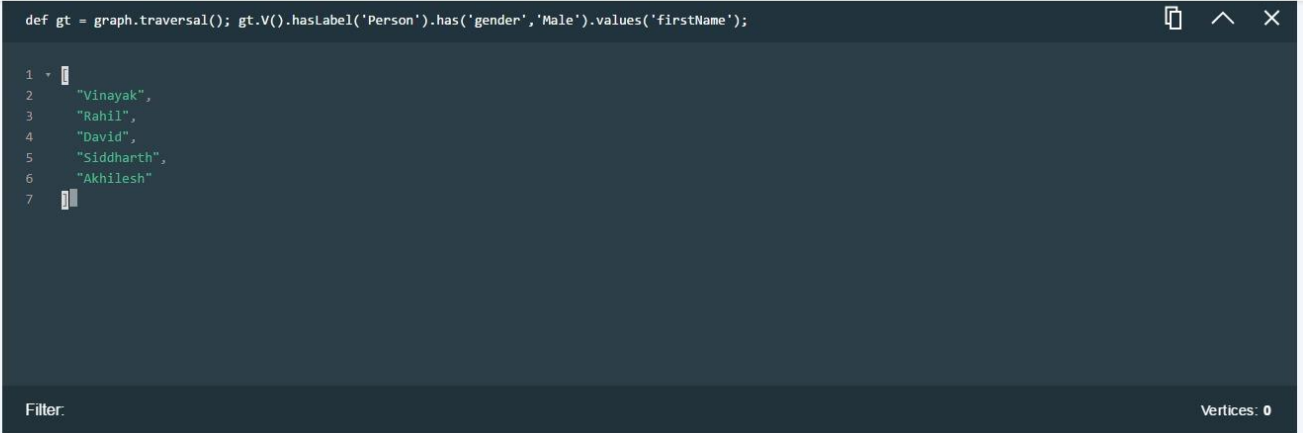
QUERY#1

Query in GREMLIN to show all the person who has gender Male

```
def gt = graph.traversal();
gt.V().hasLabel('Person').has('gender','Male').values('firstName');
```

## OUTPUT

```
1 def gt = graph.traversal(); gt.V().hasLabel('Person').has('gender','Male').values('firstName');
2
```



## QUERY#2

Query in GREMLIN to show all the details of person with firstName : Siddharth

```
def gt = graph.traversal();
gt.V().hasLabel('Person').has('firstName','Siddharth').outE('knows','isA','worksAt').inV().path();
```

OUTPUT:



## QUERY#3

Query in GREMLIN to find out all the details of people working at google with "competency level expert " of any languages .

```
def gt = graph.traversal();
gt.V().hasLabel('Company').has('companyName','Google').inE('worksAt').outV().outE('knows').has('competencyLevel','expert').inV().path();
```

OUTPUT in json :

```
[
{
  "labels": [
    [],
    []
  ]
}
```

```
[],
[],
[]
],
"objects": [
{
  "id": 4216,
  "label": "Company",
  "type": "vertex",
  "properties": {
    "established": [
      {
        "id": "1lb-394-2a6d",
        "value": "1995"
      }
    ],
    "companyName": [
      {
        "id": "173-394-sl",
        "value": "Google"
      }
    ]
  }
},
{
  "id": "e8b-388-fth-394",
  "label": "worksAt",
  "type": "edge",
  "inVLabel": "Company",
  "outVLabel": "Person",
  "inV": 4216,
  "outV": 4184,
  "properties": {
    "joiningDate": "29 Feb 2012"
  }
},
{
  "id": 4184,
  "label": "Person",
  "type": "vertex",
  "properties": {
    "firstName": [
      {
        "id": "16z-388-2dh",
        "value": "Vinayak"
      }
    ],
    "lastName": [
      {
        "id": "1l7-388-35x",
        "value": "Patel"
      }
    ],
    "address": [
      {
        "id": "2rv-388-5j9",
```

```

      "value": "San Fransisco"
    }
  ],
  "gender": [
    {
      "id": "363-388-6bp",
      "value": "Male"
    }
  ],
  "employeeId": [
    {
      "id": "3yj-388-a9x",
      "value": "G002"
    }
  ],
  "salary": [
    {
      "id": "2dn-388-4qt",
      "value": 80000
    }
  ],
  "age": [
    {
      "id": "1zf-388-3yd",
      "value": 26
    }
  ],
  "email": [
    {
      "id": "3kb-388-745",
      "value": "Vinayak@gmail.com"
    }
  ]
}
},
{
  "id": "du3-388-glx-36g",
  "label": "knows",
  "type": "edge",
  "inVLabel": "Skill",
  "outVLabel": "Person",
  "inV": 4120,
  "outV": 4184,
  "properties": {
    "competencyLevel": "expert"
  }
},
{
  "id": 4120,
  "label": "Skill",
  "type": "vertex",
  "properties": {
    "skillName": [
      {
        "id": "16r-36g-7wl",
        "value": "Javascript"
      }
    ]
  }
}

```

```

    }
  ]
}
}
]
},
{
  "labels": [
    [],
    [],
    [],
    [],
    []
  ],
  "objects": [
    {
      "id": 4216,
      "label": "Company",
      "type": "vertex",
      "properties": {
        "established": [
          {
            "id": "11b-394-2a6d",
            "value": "1995"
          }
        ],
        "companyName": [
          {
            "id": "173-394-sl",
            "value": "Google"
          }
        ]
      }
    },
    {
      "id": "emj-9js-fth-394",
      "label": "worksAt",
      "type": "edge",
      "inVLabel": "Company",
      "outVLabel": "Person",
      "inV": 4216,
      "outV": 12376,
      "properties": {
        "joiningDate": "31 Oct 2015"
      }
    },
    {
      "id": 12376,
      "label": "Person",
      "type": "vertex",
      "properties": {
        "firstName": [
          {
            "id": "93f-9js-2dh",
            "value": "Siddharth"
          }
        ]
      }
    }
  ]
}

```

```

],
"lastName": [
  {
    "id": "9hn-9js-35x",
    "value": "Mewada"
  }
],
"Email": [
  {
    "id": "bgr-9js-28lh",
    "value": "sid@gmail.com"
  }
],
"address": [
  {
    "id": "aob-9js-5j9",
    "value": "Santa Cruz"
  }
],
"Gender": [
  {
    "id": "b2j-9js-27t1",
    "value": "Male"
  }
],
"Employeeid": [
  {
    "id": "buz-9js-29dx",
    "value": "G006"
  }
],
"salary": [
  {
    "id": "aa3-9js-4qt",
    "value": 20000
  }
],
"age": [
  {
    "id": "9vv-9js-3yd",
    "value": 24
  }
]
]
},
{
  "id": "dfv-9js-glx-cnc",
  "label": "knows",
  "type": "edge",
  "inVLabel": "Skill",
  "outVLabel": "Person",
  "inV": 16392,
  "outV": 12376,
  "properties": {
    "competencyLevel": "expert"
  }
}

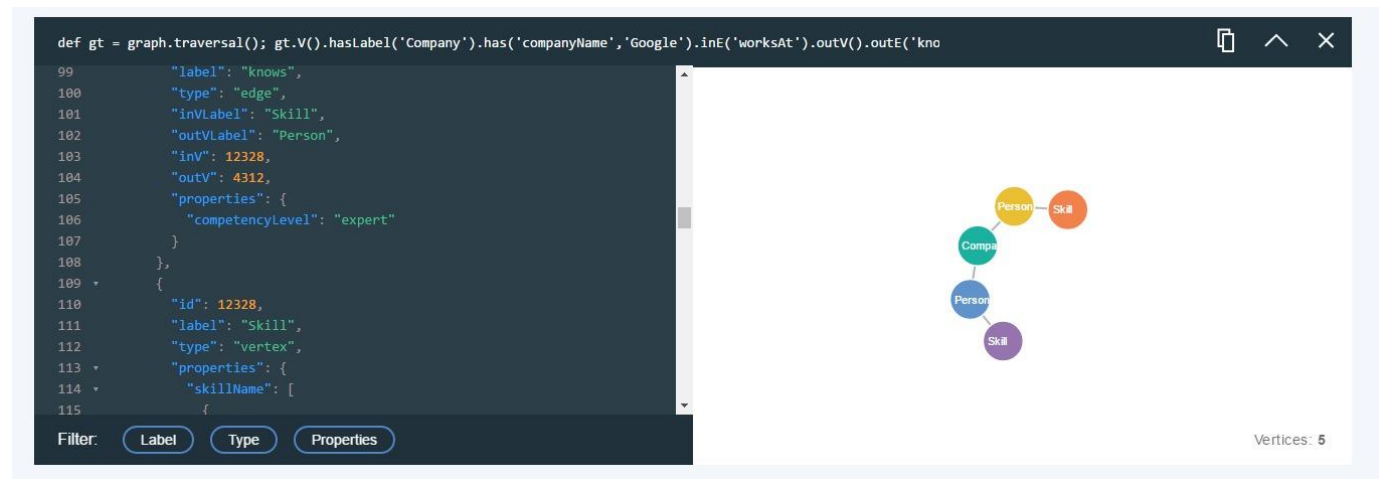
```

```

    },
    {
      "id": 16392,
      "label": "Skill",
      "type": "vertex",
      "properties": {
        "skillName": [
          {
            "id": "7i9-cnc-7wl",
            "value": "Python"
          }
        ]
      }
    }
  ]
}
]

```

OUTPUT:



## REFERENCES :

- 1) [www.w3schools.com](http://www.w3schools.com)